

# Load Balancer

Scaling up

# Objectives

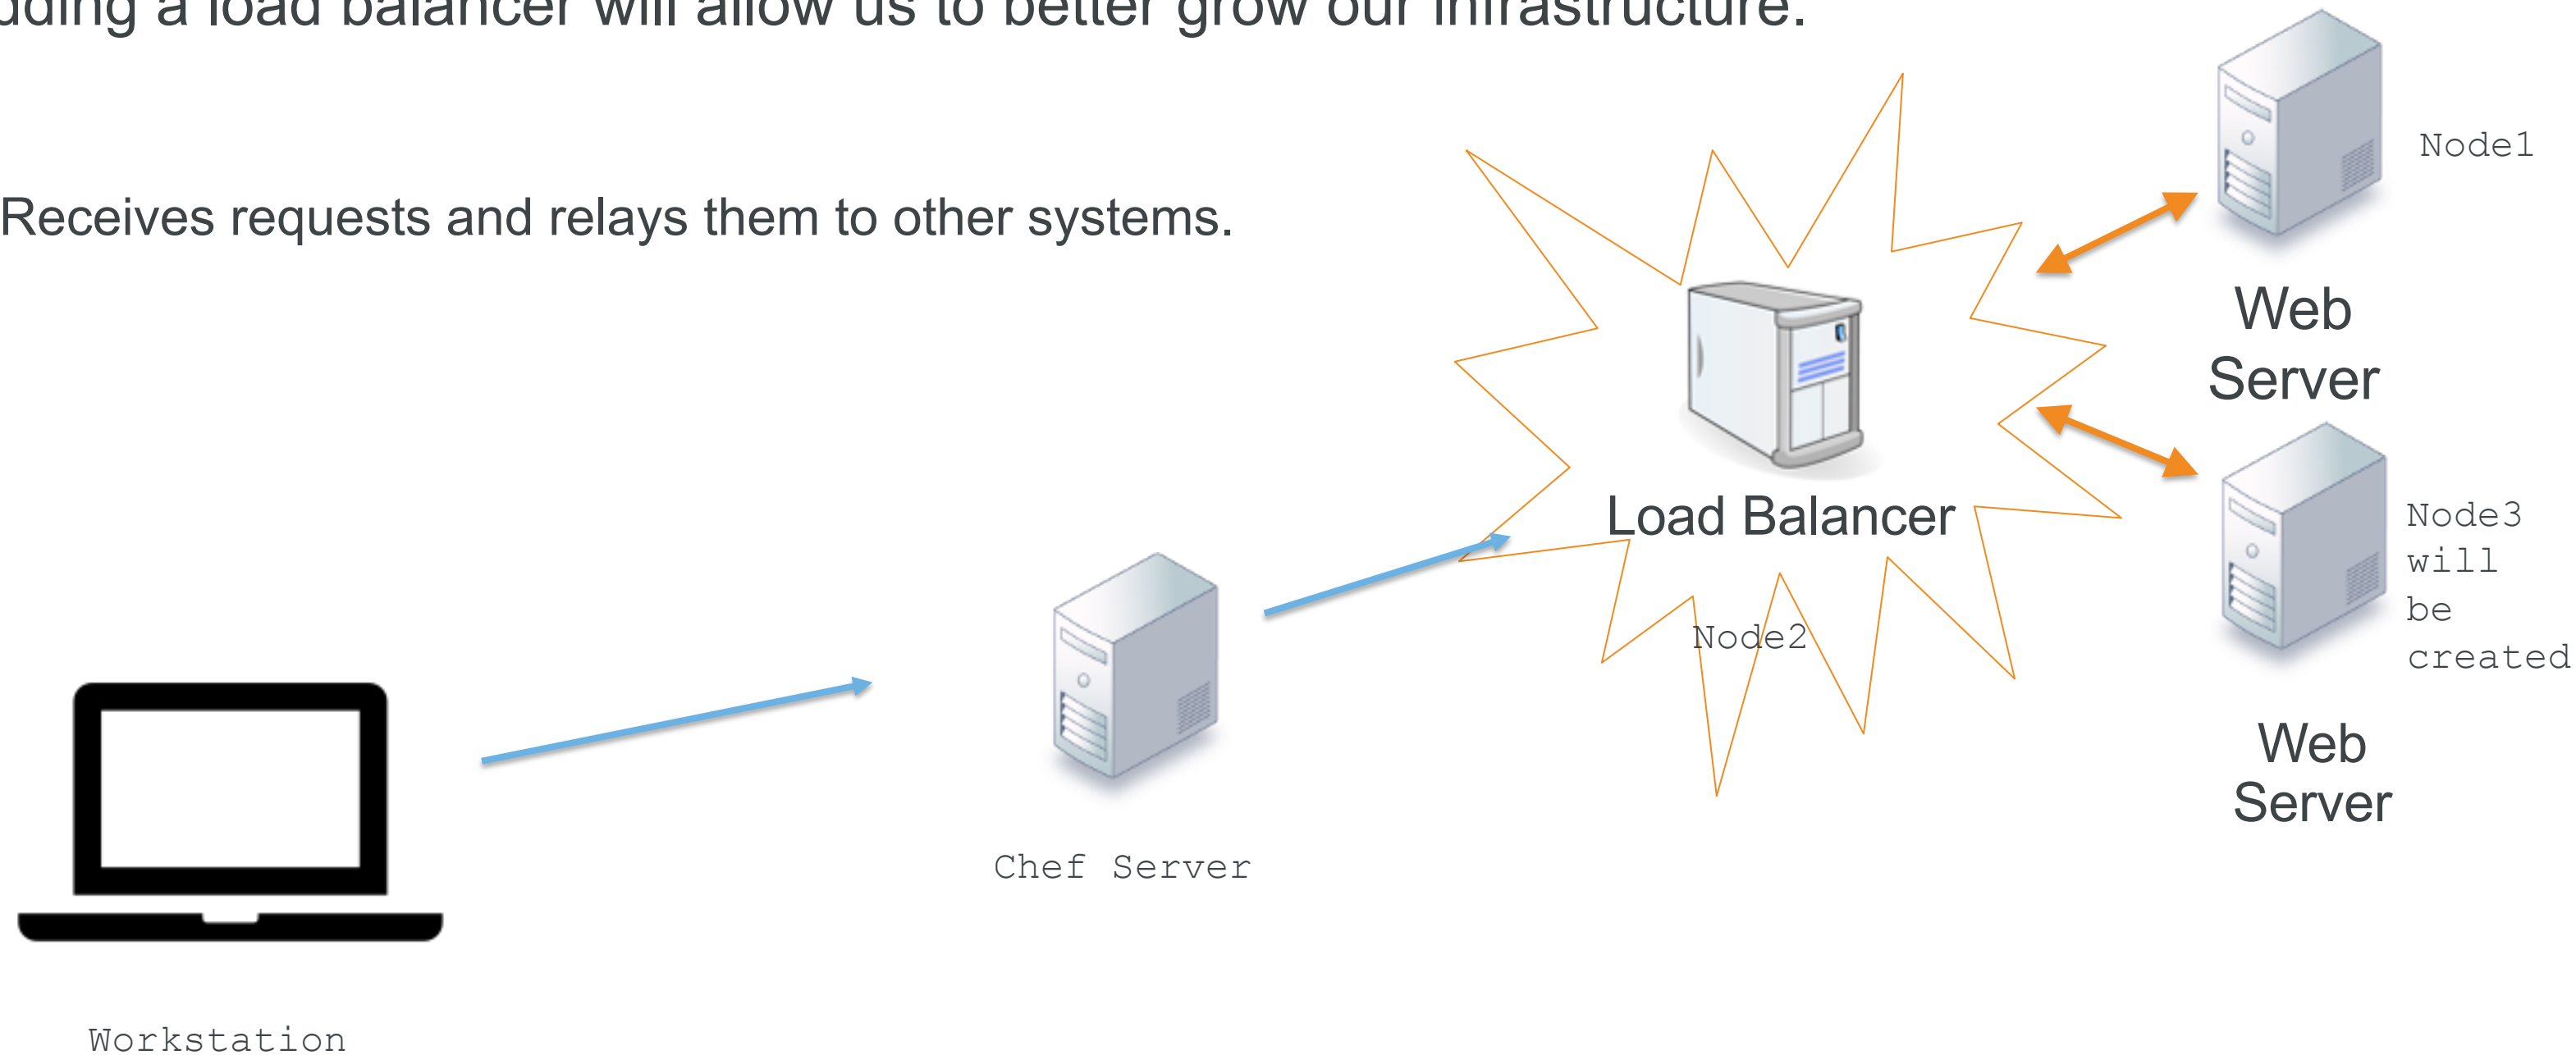
After completing this module, you should be able to

- Create a haproxy cookbook
- Add traffic route to existing webserver in node1
- Upload haproxy cookbook to chef server
- Bootstrap a node2 for haproxy

# Load Balancer

Adding a load balancer will allow us to better grow our infrastructure.

Receives requests and relays them to other systems.





# Manage Multiple Nodes

Our site has just gotten super busy so we now need to manage multiple nodes.



## GL: Scaling up

*Our site has just got super busy with multiple web servers – so we now need a load balancer.*

### Objective:

- ☐ Create a load balancer cookbook
- ☐ Upload cookbook to Chef Server
- ☐ Bootstrap a new node that runs the load balancer cookbook

# GL: Generate HAProxy Cookbook



```
$ cd ~/chef-repo
```

```
$ chef generate cookbook cookbooks/haproxy
```

```
Compiling Cookbooks...
```

```
Recipe: code_generator::cookbook
```

```
  * directory[C:/Users/YOU/chef-repo/cookbooks/haproxy] action create
    - create new directory C:/Users/YOU/chef-repo/cookbooks/haproxy
  * template[C:/Users/YOU/chef-repo/cookbooks/haproxy/metadata.rb] action create_if_missing
    - create new file C:/Users/YOU/chef-repo/cookbooks/haproxy/metadata.rb
    - update content in file C:/Users/YOU/chef-repo/cookbooks/haproxy/metadata.rb from none
to 899276
    (diff output suppressed by config)
  * template[C:/Users/YOU/chef-repo/cookbooks/haproxy/README.md] action create_if_missing
```

# GL: Edit haproxy cookbook's default recipe

 `chef-repo/cookbooks/haproxy/recipes/default.rb`

```
#  
# Cookbook Name:: haproxy  
# Recipe:: default  
#  
# Copyright (c) 2016 The Authors, All Rights Reserved.  
  
package 'haproxy'  
  
template '/etc/haproxy/haproxy.cfg' do  
  source 'haproxy.cfg.erb'  
end  
  
service 'haproxy' do  
  action [:start, :enable]  
end
```

# DISCUSSION



## Configure HAProxy

We need to configure HAProxy to route traffic to our web server.

Have a look at how HAProxy is configured

<http://bit.ly/1Xoai9R>



# GL: Generate the Template



```
$ chef generate template cookbooks/haproxy haproxy.cfg
```

```
Compiling Cookbooks...
```

```
Recipe: code_generator::template
```

```
  * directory[cookbooks/haproxy/templates/default] action create
```

```
    - create new directory cookbooks/haproxy/templates/default
```

```
  * template[cookbooks/haproxy/templates/default/haproxy.cfg.erb] action create
```

```
    - create new file cookbooks/haproxy/templates/default/haproxy.cfg.erb
```

```
    - update content in file cookbooks/haproxy/templates/default/haproxy.cfg.erb  
from none to e3b0c4
```

```
    (diff output suppressed by config)
```

11/09/2016

# GL: Configuring haproxy.cfg.erb

 cookbooks/haproxy/templates/haproxy.cfg.erb

```
...
frontend main *:80
  acl url_static      path_beg      -i /static /images /javascript /stylesheets
  acl url_static      path_end      -i .jpg .gif .png .css .js

  use_backend static   if url_static
  default_backend      app

backend static
  balance roundrobin
  server static 127.0.0.1:4331 check

backend app
  balance roundrobin
  server app <<IP ADDRESS>>:80 weight 1 maxconn 100 check
```

Copy the bit.ly URL from below the slide in your participant guide:

<http://bit.ly/1Xoai9R>

# DISCUSSION



## Deriving the Public IP Address in AWS

Let's stop for a moment and discuss how we can add a webserver's public IP Addresses to haproxy.cfg using the `cloud` attribute.

# HAProxy Configuration Should Look Like This

```
...  
backend app  
    balance roundrobin  
    server app0 <<node1 IP ADDRESS>>:80 weight 1 maxconn 100 check  
    server app1 <<node2 IP ADDRESS>>:80 weight 1 maxconn 100 check
```

We need to add the webserver's IP Addresses to haproxy.cfg

Doing it manually seems wrong

Should be able to Search all desired webserver IPs !!!

# HAProxy Configuration Should Look Like This

```
...  
backend app  
  balance roundrobin  
  server app0 <<node1 IP ADDRESS>>:80 weight 1 maxconn 100 check  
  server app1 <<node2 IP ADDRESS>>:80 weight 1 maxconn 100 check
```



We will do search capability in haproxy later, let's hard code now and test haproxy cookbook

**But we need IP, you can use node1 IP, but there is a better way**

# DISCUSSION



## Introducing Chef Server SOLR search

Chef server use SOLR for indexing and searching.  
Node objects, attributes, cookbook details, recipe details –  
all gets indexed in SOLR and can be searched.  
To search IP, we will use “knife search node”

# Knife search for IPs



```
$ knife search node "policy_group:prod AND policy_name:company_web"
```

```
1 items found
```

```
Node Name:    node1
```

```
Policy Name:  company_web
```

```
Policy Group: prod
```

```
FQDN:         ip-172-31-18-33.us-east-2.compute.internal
```

```
IP:           3.16.131.207
```

```
Run List:     recipe[company_web::default]
```

```
Recipes:      company_web::default, apache::default, apache::server
```

```
Platform:     centos 7.8.2003
```

```
Tags:
```

# DISCUSSION



## Amazon EC2 Instances

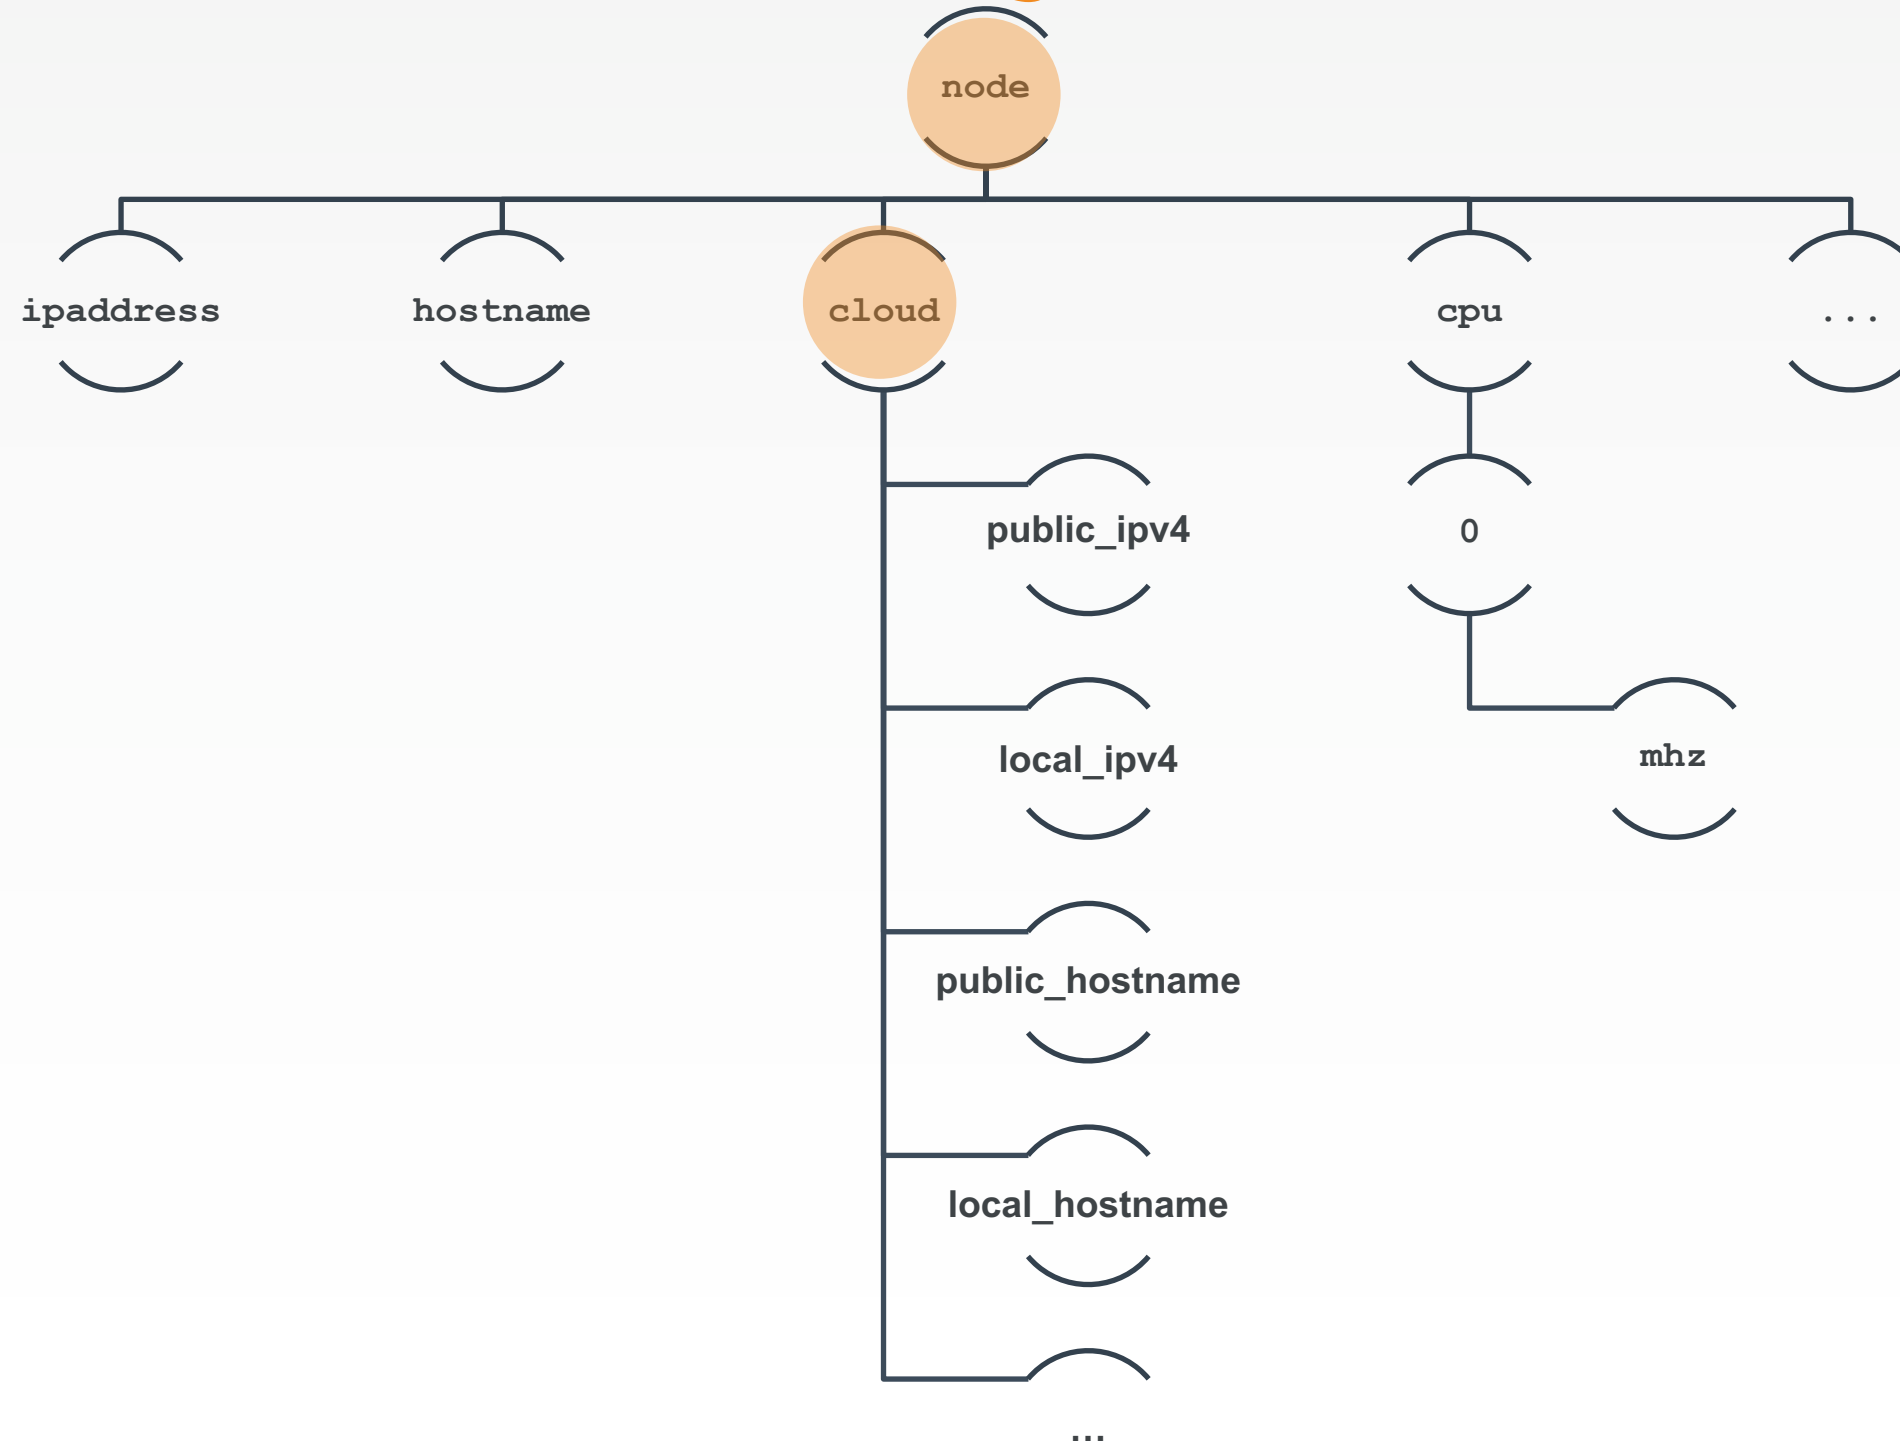
We can't use the `ipaddress` attribute within our recipes – they're on the private (internal) network and we need external access.

In a previous section we looked at the node object.

Nodes in EC2 have some specific networking attributes in their node object – some private & some public.



# EC2 and the Node Object



# Individual Node's EC2 Information



```
$ knife node show node1 -a cloud
```

```
node1:
```

```
cloud:
```

```
local_hostname: ip-172-31-29-218.ec2.internal
```

```
local_ipv4: 172.31.29.218
```

```
private_ips: 172.31.29.218
```

```
provider: ec2
```

```
public_hostname: ec2-54-88-185-159.compute-1.amazonaws.com
```

```
public_ips: 54.88.185.159
```

```
public_ipv4: 54.88.185.159
```

LOCAL

# Individual Node's EC2 Information



```
$ knife node show node1 -a cloud.public_ipv4
```

```
node1:
```

```
cloud.public_ipv4: 3.16.131.207
```

LOCAL

# Individual Node's Public IP Address



```
$ knife search node "policy_group:prod AND  
policy_name:company_web" -a cloud.public_ipv4
```

```
1 items found
```

```
node1:
```

```
cloud.public_ipv4: 3.16.131.207
```

LOCAL

# HAProxy Configuration Should Look Like This

```
...  
backend app  
    balance roundrobin  
    server app 3.16.131.207:80 weight 1 maxconn 100 check
```



## Scaling up


*Our site has just got super busy with multiple web servers – so we now need a load balancer.*

- ✓ Create a load balancer cookbook
- ☐ Upload cookbook to Chef Server
- ☐ Bootstrap a new node that runs the load balancer cookbook

# Upload the Policyfile



```
$ cd ~/chef-repo
```



LOCAL

# Generate Policyfile



```
$ chef generate policyfile haproxy
```

```
Resolving cookbook dependencies...
```

```
Fetching 'haproxy' from source at .
```

```
Fetching cookbook index from https://supermarket.chef.io...
```

```
Using haproxy (0.1.0) from source at .
```

LOCAL



# Change Policyfile content



```
$ vi ~/chef-repo/haproxy.rb
```

```
# A name that describes what the system you're building with Chef does.  
name 'haproxy'  
  
# Where to find external cookbooks:  
default_source :supermarket  
  
# run_list: chef-client will run these recipes in the order specified.  
run_list 'haproxy::default'  
  
# Specify a custom source for a single cookbook:  
cookbook 'haproxy', path: 'cookbooks/haproxy'
```

# Install Policyfile



```
$ chef install haproxy.rb
```

```
Building policy haproxy
```

```
Expanded run list: recipe[haproxy::default]
```

```
Caching Cookbooks...
```

```
Installing haproxy >= 0.0.0 from path
```

```
Lockfile written to /home/centos/chef-repo/haproxy.lock.json
```

```
Policy revision id: 648817382c99f23416f7ff2a3d14da7e812bf6d184bf6e4f810ed49084ba6de1
```

LOCAL

# Push Policyfile to Server



```
$ chef push prod haproxy.lock.json
```

```
Uploading policy haproxy (648817382c) to policy group prod  
Uploaded haproxy 0.1.0 (d699a0c0)
```

LOCAL

# Verify the Policyfile Upload



```
$ chef show-policy
```

```
apache
```

```
* prod:  ab5d900157
```

```
company_web
```

```
* prod:  c5590c5d36
```

```
haproxy
```

```
* prod:  648817382c
```

```
workstation
```

```
* prod:  2e5817026e
```



## Lab: Scaling up

*Our site has just got super busy with multiple web servers – so we now need a load balancer.*

### Objective:

- ✓ Create a load balancer cookbook
- ✓ Upload cookbook to Chef Server
- ❑ Bootstrap a new node that runs the load balancer cookbook

# Bootstrap a Load Balancer Node



```
$ knife bootstrap <ip_node2> -x centos --sudo -i <aws.pem path> -N node2 --  
policy-name haproxy --policy-group prod
```

This will bootstrap a node, and also apply haproxy policy

LOCAL

# Validate the New Node



```
$ knife node show node2
```

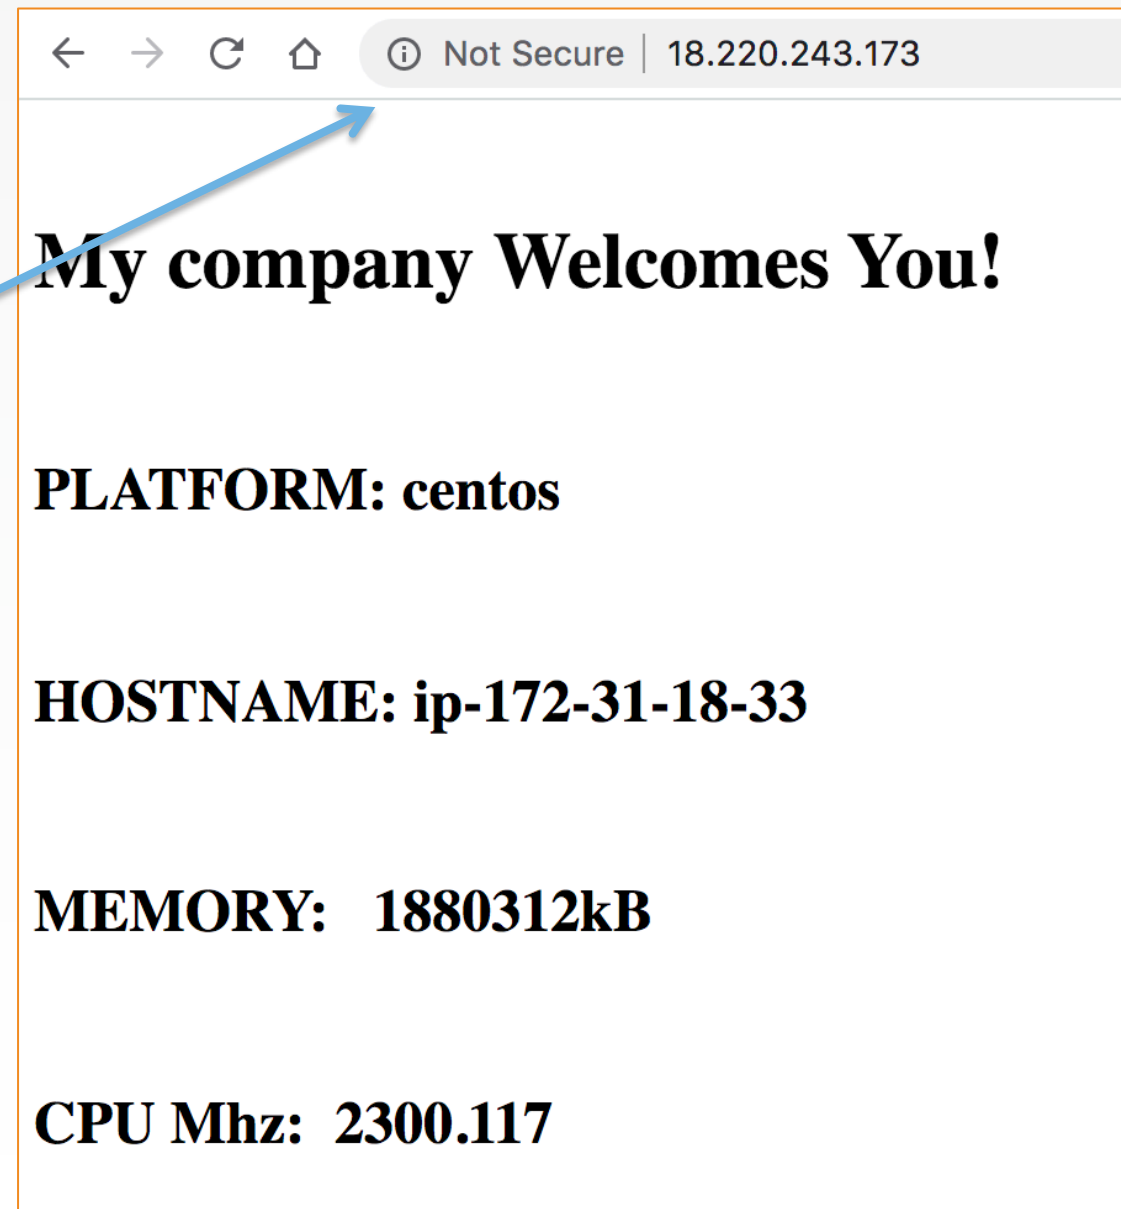
```
Node Name:    node2
Policy Name:   haproxy
Policy Group:  prod
FQDN:         ip-172-31-19-183.us-east-2.compute.internal
IP:           18.220.243.173
Run List:      recipe[haproxy::default]
Recipes:       haproxy::default
Platform:      centos 7.8.2003
Tags:
```

LOCAL

# Validate Load Balancer

Load balancer  
IP

Routing to webserver







# Getting Second Webserver

## Objective:

- Bootstrap third node – our second webserver
- Manually change loadbalancer
- Upload to server
- Test Load balancer

# Bootstrap a web server Node



```
$ knife bootstrap <ip_node3> -x centos --sudo -i <aws.pem path> -N node3 --  
policy-name company_web --policy-group prod
```

This will bootstrap a node, and also apply haproxy policy

LOCAL

# Validate the New Node



```
$ knife node show node3
```

```
Node Name:    node3
Policy Name:   company_web
Policy Group:  prod
FQDN:         ip-172-31-24-5.us-east-2.compute.internal
IP:           52.15.221.52
Run List:     recipe[company_web::default]
Recipes:      company_web::default, apache::default, apache::server
Platform:     centos 7.8.2003
Tags:
```

Use this IP  
Add to loadbalancer

LOCAL



# Upgrading Load Balancer

## Objective:

- Manually change loadbalancer
- Upload to server
- Test Load balancer

# HAProxy Configuration Should Look Like This

```
$ ~/chef-repo/cookbooks/haproxy/templates/haproxy.cfg.erb
```

```
...  
backend app  
  balance roundrobin  
  server app0 3.16.131.207:80 weight 1 maxconn 100 check  
  server app1 52.15.221.52:80 weight 1 maxconn 100 check
```

# Upload the metadata.rb



```
$ ~/chef-repo/cookbooks/haproxy/metadata.rb
```

```
name 'haproxy'  
maintainer 'The Authors'  
maintainer_email 'you@example.com'  
license 'All Rights Reserved'  
description 'Installs/Configures haproxy'  
long_description 'Installs/Configures haproxy'  
version '0.2.0'  
chef_version '>= 14.0'
```

Change version, very important !!

LOCAL

# Go to chef-repo



```
$ cd ~/chef-repo
```

LOCAL

# Update Policyfile



```
$ chef update haproxy.rb
```

```
Attributes already up to date
```

```
Building policy haproxy
```

```
Expanded run list: recipe[haproxy::default]
```

```
Caching Cookbooks...
```

```
Installing haproxy >= 0.0.0 from path
```

```
Lockfile written to /home/centos/chef-repo/haproxy.lock.json
```

```
Policy revision id:
```

```
f501a623a03dec00478641e67ac60a72ccc813888daf22e1a42783ff4182e061
```

LOCAL




# Push Policyfile to Chef Server



```
$ chef push prod haproxy.lock.json
```

```
Uploading policy haproxy (f501a623a0) to policy group prod  
Uploaded haproxy 0.2.0 (ae0efc25)
```



Updated version now pushed to the chef server

LOCAL

# Apply new changes to LB node



```
$ knife ssh <ip_load balancer> -m -x centos -i  
<aws.pem path> "sudo chef-client"
```

```
Building policy haproxy
```

```
Expanded run list: recipe[haproxy::default]
```

```
Caching Cookbooks...
```

```
Installing haproxy >= 0.0.0 from path
```

```
Lockfile written to /home/centos/chef-repo/haproxy.lock.json
```

```
Policy revision id: 648817382c99f23416f7ff2a3d14da7e812bf6d184bf6e4f810ed49084ba6de1
```

LOCAL

# Lab: Test the Load Balancer

## Note:

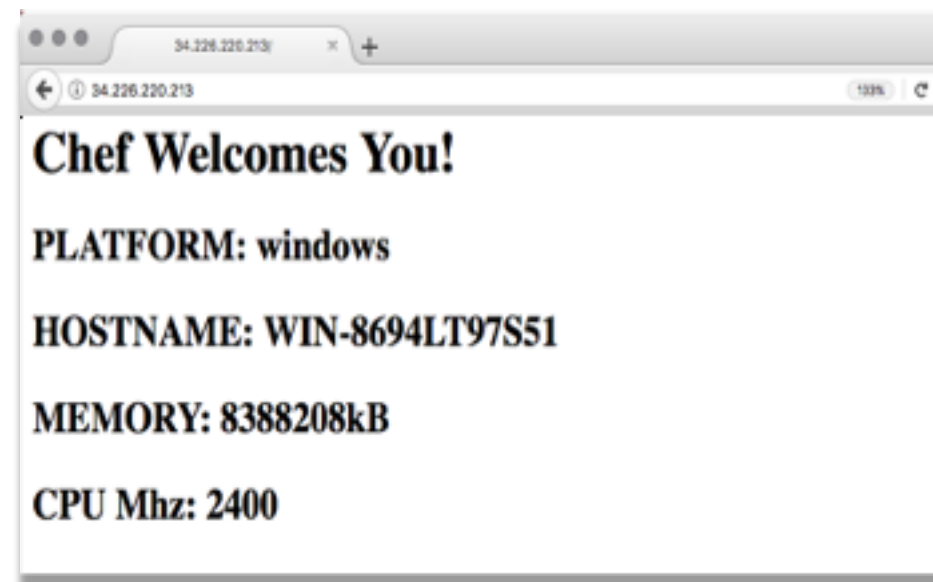

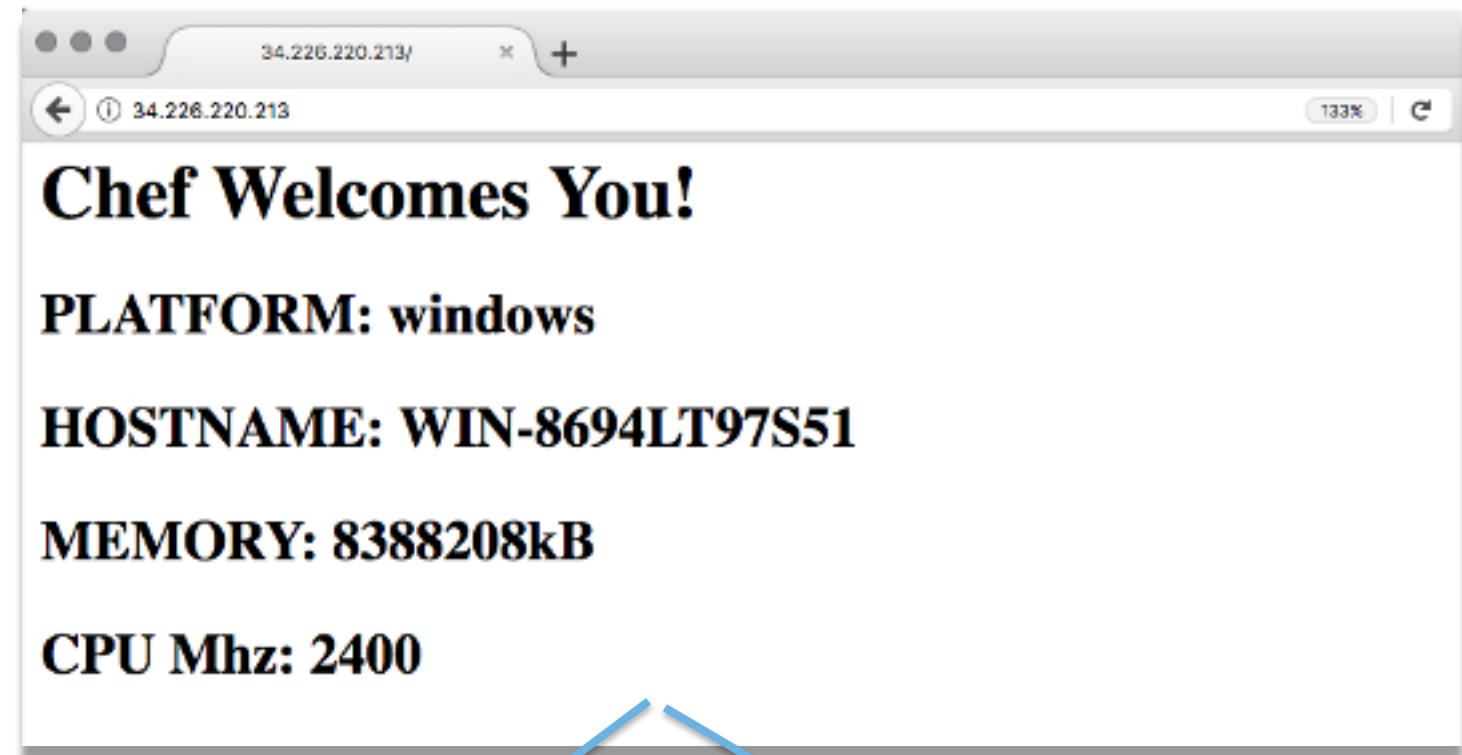
Haproxy reads from cache.

Thus, any change in config needs a restart of the service.

You can ssh to the LB node and run “sudo systemctl restart haproxy” to restart the haproxy service

-- We will do this in a much elegant way in the next chapter, as it needs some background of **knife search**

LB



Server 1



Server 2



**CHEF**™